

HENRY MOSS

STOR-i, Science and Technology Building, Lancaster University, LA1 4YR
07837 072677 ◊ h.moss@lancaster.ac.uk

RESEARCH INTERESTS

Bayesian optimisation: information-theory; batch design; multi-fidelity.

Gaussian processes: multi-fidelity models; string kernels.

Experimental design: gene design; molecule search.

Natural language processing: AutoML; hyper-parameter tuning; text-to-speech.

CURRENT ROLE

Senior Machine Learning Researcher, Secondmind, Cambridge.

EDUCATION

PhD in Machine Learning and Statistics, STOR-i CDT, Lancaster University. 2016 - 2017

- *Information-theoretic Bayesian optimisation for natural language processing:*
 - Developing Bayesian optimisation methods for high-cost string design problems.
 - Extending information-theoretical Bayesian optimisation for batch and multi-fidelity designs.
- Supervision by Prof. David Leslie (Statistics) and Prof. Paul Rayson (Computer Science).

MRes in Statistics and Operational Research (Distinction), Lancaster University. 2016 - 2017

MA (Hons) in Mathematics (2:1), Emmanuel College, University of Cambridge. 2013 - 2016

PUBLICATIONS

Published:

Vakili S., Moss H. B., Artmev A., Dutordoir V. & Picheny V. Scalable Thompson Sampling using Sparse Gaussian Process Models. In *The Conference on Neural Information Processing Systems* (NeurIPS), 2021.

Moss H. B., Leslie D. S., Gonzalez J. & Rayson P. General-purpose Information-based Bayesian Optimisation. In *The Journal for Machine Learning Research* (JMLR), 2021.

Moss H. B., Beck D., Leslie D. S., Gonzalez J. & Rayson P. Bayesian Optimisation over String spaces. In *The Conference on Neural Information Processing Systems* (NeurIPS), 2020 (**spotlight**).

Moss H. B. & Griffiths R. Gaussian Process Molecule Property Prediction With FlowMO. In *The Conference on Neural Information Processing Systems: Machine Learning for Molecules Workshop* (NeurIPS: ML4Molecules Workshop), 2020.

Moss H. B., Leslie D. S. & Rayson P. BOSH: Bayesian Optimisation by Sampling Hierarchically. In *The International Conference on Machine Learning: Workshop on Real World Experimental Design and Active Learning* (ICML: Real-ML Workshop), 2020.

Moss H. B., Leslie D. S. & Rayson P. MUMBO: Multi-task Max-value Bayesian Optimisation. In *The European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases* (ECML), 2020.

Moss H. B., Aggarwal V., Prateek N., Gonzalez J. & Barra-Chicote R. BOFFIN TTS: Few-shot Speaker Adaptation By Bayesian Optimisation. In *The International Conference on Acoustics, Speech and Signal Processing* (ICASSP), 2020.

Moss H. B., Moore A., Leslie D. S. & Rayson P. FIESTA: Fast IdEntification of State-of-The-Art Models Using Adaptive Bandit Algorithms. In *The Annual Meeting of the Association of Computational Linguists* (ACL), 2019.

Moss H. B., Leslie D. S. & Rayson P. Using *J-K*-fold Cross Validation to Reduce Variance when Tuning Natural Language Processing Models. In *The International Conference on Computational Linguistics* (COLING), 2018 (**area chair favourite**).

Under Review:

Griffiths R., Thawani A., Jamasb A., **Moss H. B.**, Bourached A., Jones P., McCorkindale W., & Aldrick A. A Case for Domain Expert Dataset Curation in Machine-Learning Enabled Chemistry.

PRIZES

NeurIPS spotlight paper: Top 3% of submissions at NeurIPS.	2020
ML4Molecules contributed talk: Top 5% of submissions at the ML4Molecules Workshop at NeurIPS.	2020
Nick Smith prize: Best second-year Statistics PhD student at Lancaster University.	2019
Area chair favourite: Nominated for overall best paper at COLING 2018.	2018

INTERNSHIPS AND SUMMER SCHOOLS

Visiting researcher , School of Computing and Information Systems, University of Melbourne. <ul style="list-style-type: none">Derived Bayesian optimisation for sequence design under syntactic constraints.	2020
Amazon PhD internship , Text-To-Speech Team, Amazon Alexa, Cambridge. <ul style="list-style-type: none">Used Bayesian optimisation to fine-tune neural systems to synthesise new voices with limited data.	2019
Amazon Intern Colloquium , Amazon Research, Cambridge.	2019
Google NLP Summit , Google Research, Zurich.	2019
Microsoft AI Summer School , Microsoft Research, Cambridge.	2018
UCREL NLP Summer School , Lancaster University.	2017
Wellcome Sanger internship , University of Cambridge. <ul style="list-style-type: none">Designed a system to automatically flag promising compounds during image-based drug screening.	2016
Summer research internship , STOR-i CDT, Lancaster University.	2016
Equity analyst: Oil & Gas researcher at Redburn International.	2015

PRESENTATIONS

Gaussian Process Summer School: <i>Overview of Secondmind's Toolboxes</i> (talk).	2021
ECML: <i>MUMBO: Multi-task Max-value Bayesian Optimisation</i> (talk).	2020
Microsoft Research Cambridge: <i>Bayesian Optimisation in Gene Design Loops</i> (talk).	2020
Mathematics of Data Science Conference, ICML: RealML Workshop: <i>BOSH</i> (talk).	2020
ICASSP, Amazon Research Cambridge and University of Melbourne: <i>BOFFIN TTS</i> (talk).	2020
Prowler.io, Sheffield, Lancaster and Manchester Universities; <i>MUMBO</i> (talk).	2019
Amazon Intern Colloquium: <i>Rapid Speaker Adaptation with Bayesian Optimisation</i> (poster).	2019
ACL: <i>FIESTA: Fast Identification of SOTA</i> (talk).	2019
Google NLP Summit: <i>Reliable and Efficient Hyper-parameter Tuning for NLP</i> (poster).	2019
STOR-i Forum and Lancaster Data Science Group: <i>A Crash Course in Bayes Opt</i> (talk).	2018
Microsoft AI Summer School and COLING: <i>Using J-K-fold Cross Validation ...</i> (poster)	2018
UCREL Summer School in Corpus-based NLP: <i>Instabilities in NLP models</i> (poster).	2017

RESEARCH SUPPORT

Visiting researcher grant: Collaboration sponsorship from the University of Melbourne.	2020
STOR-i research fund: Support for visit to the University of Melbourne from Lancaster University.	2020
Workshop sponsorship: Support for Bayesian optimisation workshop from Amazon Research.	2019
Faculty of Science and Technology travel grant: Travel support from Lancaster University.	2018
STOR-i PhD scholarship: Full funding for MRes and PhD programme.	2016

OTHER RELEVANT EXPERIENCE

Program Committee: Committee member for the NeurIPS ML4Molecules Workshop.	2020
Reviewer: Reviewed manuscripts for AISTATS and ECML .	2020
MSc masterclass: Designed and taught short course on Bayesian optimisation.	2019 - 2020
Intern supervisor: Supervised undergraduate statistics project, Lancaster University.	2018
Statistical consultant, Natural Language Identification Project, Lancaster University.	2018
STOR-i computing team: Assist peers with coding and distributed computing.	2017 - 2019
Tutor: Undergraduate Mathematics and MSc Data Science.	2017 - 2019
Lancaster University outreach: Interactive sessions with local secondary schools.	2016 - 2018

COMPUTING SKILLS

Python, Tensorflow, PyTorch, GPflow, Git, Maintainer of Trieste, Contributor to Emukit, GPy, GPyOpt, Cython, MXNet, R, C